

Artificial Intelligence

CS-401



Chapter # 03

Solving Problems by Searching

Dr. Hafeez Ur Rehman

(Email: hafeez.urrehman@nu.edu.pk)

Chapter's Outline

- ❑ Problem Solving Agents
- ❑ Well Defined Problems and Solutions
- ❑ Example Problems
- ❑ Searching for Solutions
 - Infrastructure of Search Algorithms
 - Performance Metric
- ❑ Uninformed Search Strategies
- ❑ Informed (Heuristic) Search Strategies

Recall: Types of agents

Reflex agent



- Consider how the world IS
- Choose action based on current percept
- Do not consider the future consequences of actions

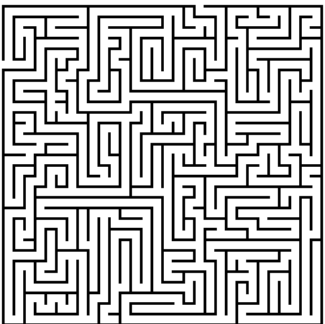
Planning agent



- Consider how the world WOULD BE
- Decisions based on (hypothesized) consequences of actions
- Must have a model of how the world evolves in response to actions
- Must formulate a goal

Initial Simplifying Assumptions

- Environment is **static**
 - no changes in environment while problem is being solved
- Environment is **observable**
- Environment and actions **are discrete**
 - (typically assumed, but we will see some exceptions)
- Environment is **deterministic**
- Environment is **known**




8			4		6			7
	1					4		
						6	5	
5		9		3		7	8	
				7				
	4	8		2		1		3
	5	2						9
		1						
3			9		2			5



Search

- We will consider the problem of designing **goal-based agents** in **fully observable, deterministic, discrete, known** environments

Start state



Goal state

Search

- We will consider the problem of designing **goal-based agents** in **fully observable, deterministic, discrete, known** environments
 - The agent must find a *sequence of actions* that reaches the goal
 - The **performance measure** is defined by
 - A. Reaching the goal and
 - B. How “**expensive**” the path to the goal is
 - We are focused on the process of finding the solution; while executing the solution, we assume that the agent can safely ignore (through abstraction) its percepts.

Abstraction

- **Definition:**

Process of removing irrelevant detail to create an abstract representation: ``high-level'', ignores irrelevant details

- **Navigation Example: how do we define states and operators?**

- First step is to abstract “the big picture”

- i.e., solve a map problem
- nodes = cities, links = freeways/roads (a high-level description)
- this description is an abstraction of the real problem


- Can later worry about details like scenery, refueling, etc

- **Abstraction is critical for automated problem solving**

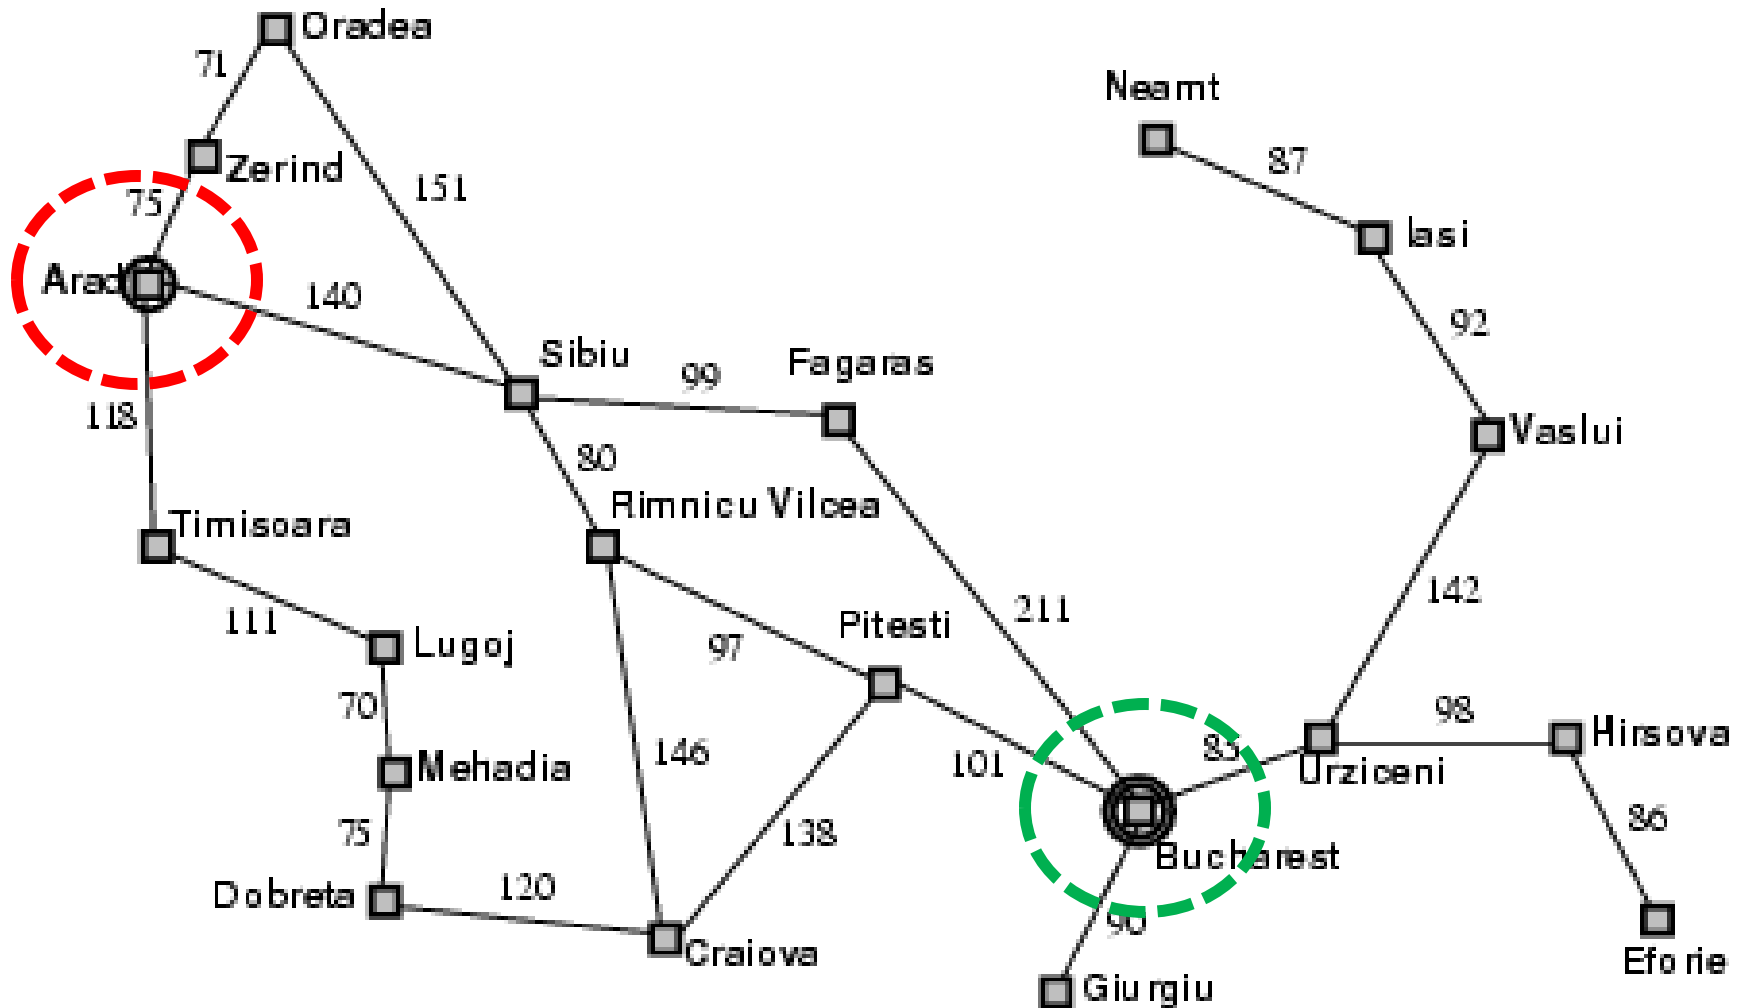
- must create an approximate, simplified, model of the world for the computer to deal with: real-world is too detailed to model exactly
- **good abstractions retain all important details**

Search problem components

- **Initial state**
 - **Actions**
 - **Transition model**
 - What state results from performing a given action in a given state?
 - **Goal state**
 - **Path cost**
 - Assume that it is a sum of nonnegative *step costs*
- The **optimal solution** is the sequence of actions that gives the *lowest* path cost for reaching the goal

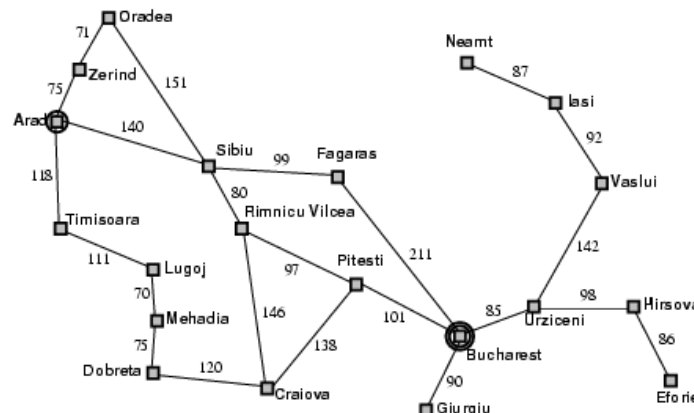


Example: Traveling in Romania




State space

- The initial state, actions, and transition model define the **state space** of the problem
 - The set of all reachable states from an initial state by any sequence of actions is called the state space
 - Can be represented as a **directed graph** where the nodes are states and links between nodes are actions
- What is the state space for the Romania problem?




Example 1 : Traveling Salesperson Problem

- Find the shortest tour that visits all cities without visiting any city twice and return to starting point.
- State: sequence of cities visited
- $S_0 = A$



- $G =$ a complete tour


Example 2: Vacuum world



- **States**
 - Agent location and dirt location
 - How many possible states?
 - What if there are n possible locations?
 - The size of the state space grows exponentially with the “size” of the world!
- **Actions**
 - Left, right, suck

Vacuum world state space graph

- Transition model



Example 3: The 8-puzzle

- **States**

- Locations of tiles

- 8-puzzle: 181,440 states ($9!/2$)
 - 15-puzzle: ~10 trillion states
 - 24-puzzle: $\sim 10^{25}$ states

- **Actions**

- Move blank left, right, up, down

- **Path cost**

- 1 per move

7	2	4
5		6
8	3	1

Start State



	1	2
3	4	5
6	7	8

Goal State

Example 4: 8-queens problem


Problem

Solution



State-Space problem formulation

- states?
 - any arrangement of $n \leq 8$ queens
 - or arrangements of $n \leq 8$ queens in leftmost n columns, 1 per column, such that no queen attacks any other.
- initial state? no queens on the board
- actions?
 - add queen to any empty square
 - or add queen to leftmost empty square such that it is not attacked by other queens.
- goal test? 8 queens on the board, none attacked.
- path cost? 1 per move



Search


- Given:
 - Initial state
 - Actions
 - Transition model
 - Goal state
 - Path cost
- How do we find the **optimal solution**?
 - How about building the state space and then using Dijkstra's shortest path algorithm?
 - Complexity of Dijkstra's is $O(E + V \log V)$, where V is the size of the state space
 - For AI problems in particular the state space may be huge!

Search: Basic idea

- Let's begin at the start state and **expand** it by making a list of all possible successor states
- We NEVER build the ENTIRE state space at once
- Maintain a **frontier** or a **list of unexpanded states**
- At each step, pick a state from the frontier to expand
- Keep going until you reach a **goal** state
- Try to expand as few states as possible

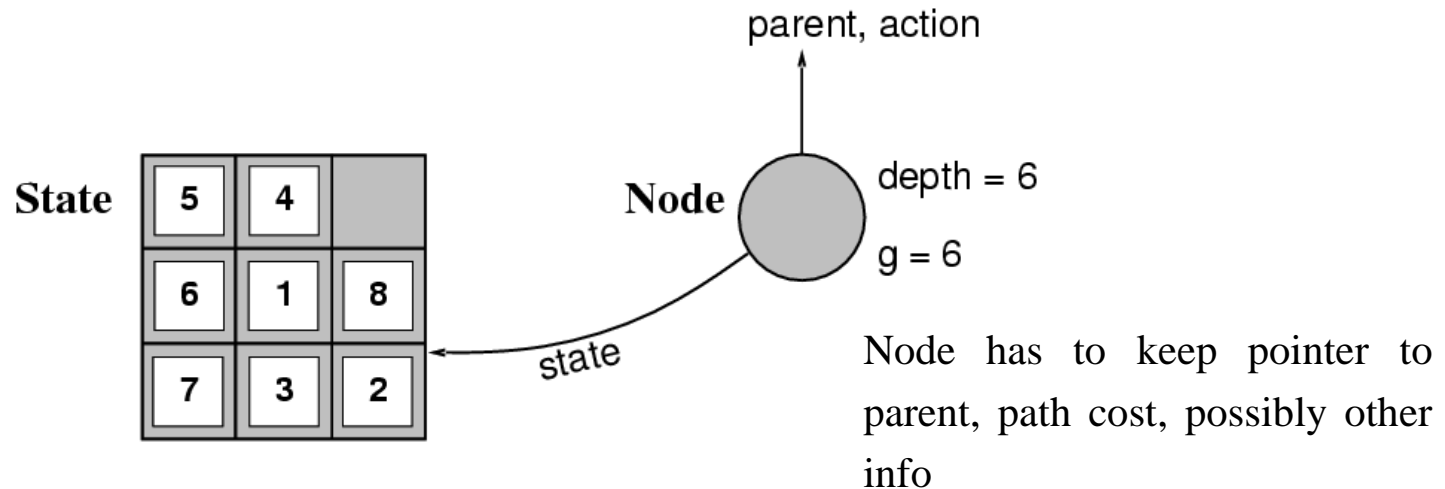
Search Graph to Tree

- “What if” tree of sequences of actions and outcomes
- The root node corresponds to the starting state
- The children of a node correspond to the **successor states** of that node’s state
- A path through the tree corresponds to a sequence of actions
 - A solution is a path ending in the goal state



States versus Nodes

- A **state** is a (representation of) a physical configuration
- A **node** is a data structure constituting part of a search tree contains info such as: **state**, **parent node**, **action**, **path cost $g(x)$** , **depth**



- The `Expand` function creates new nodes, filling in the various fields and using the `SuccessorFn` of the problem to create the corresponding states.

Search Tree for the 8 puzzle problem

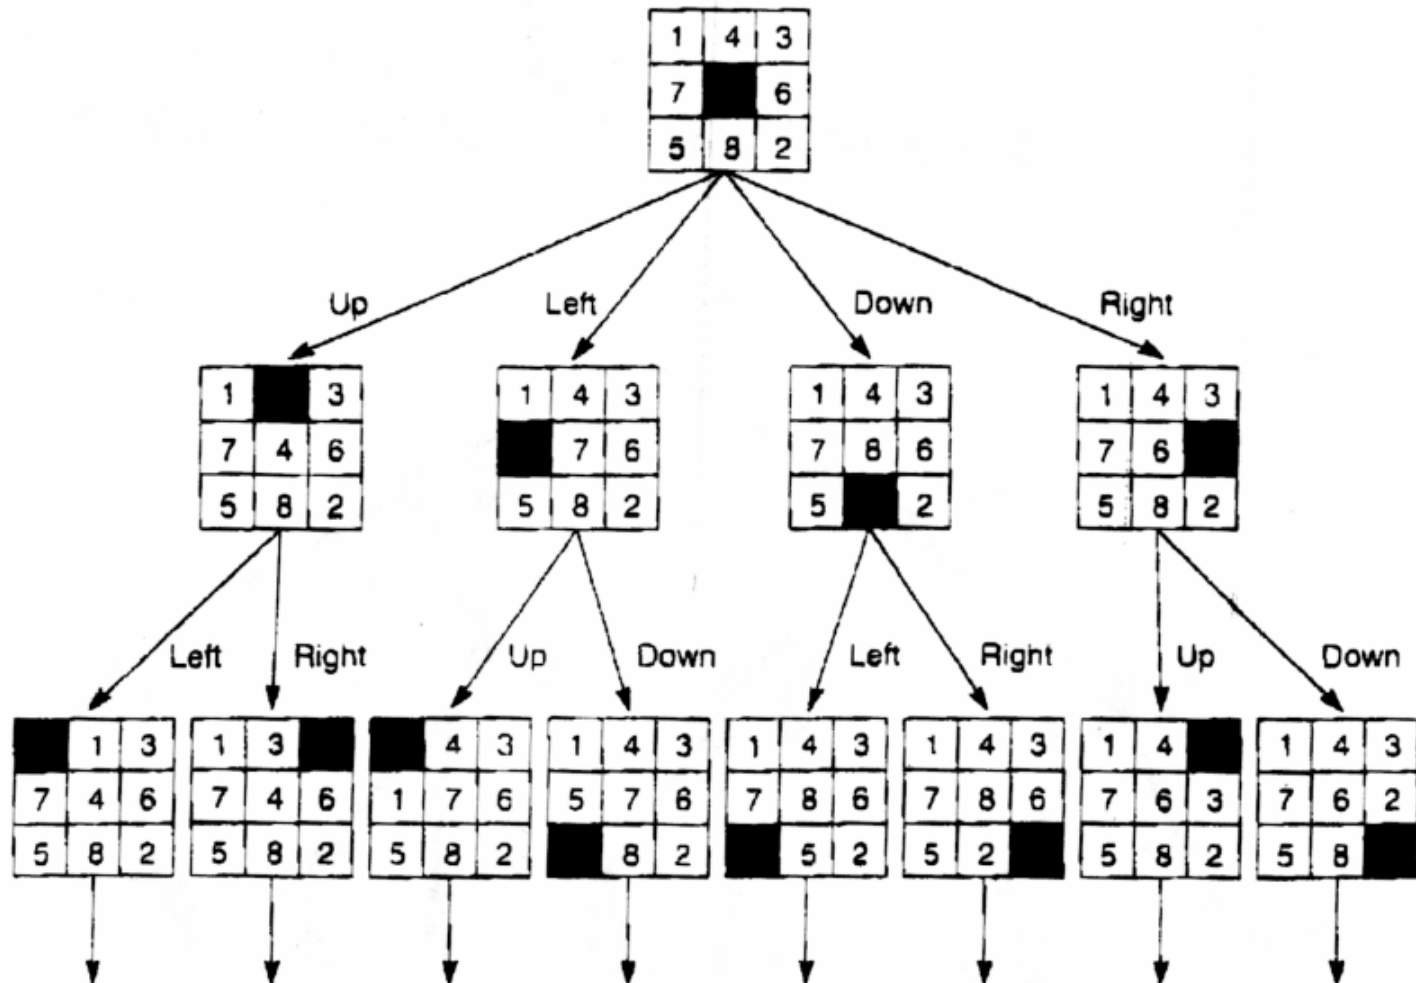



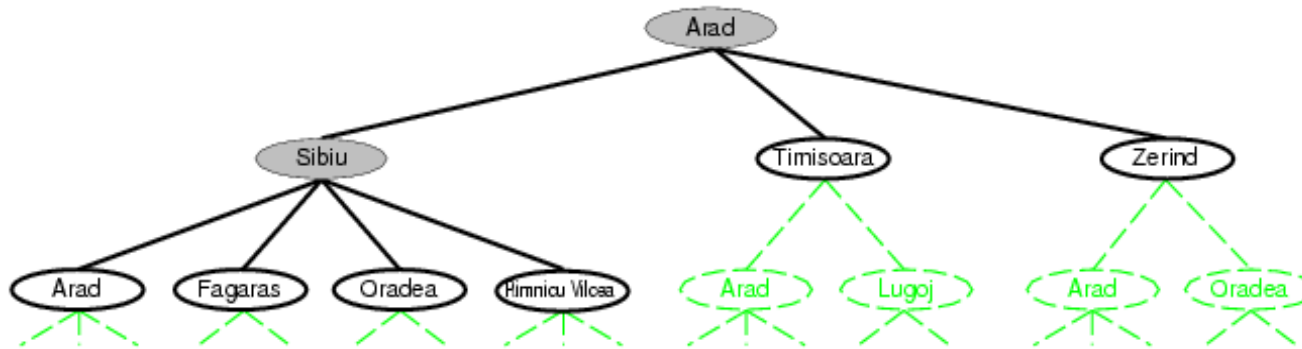
Figure 3.6 State space of the 8-puzzle generated by "move blank" operations.

Tree Search Algorithm



```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
```

Tree Search Algorithm



This “strategy” is what differentiates different search algorithms

```
function TREE-SEARCH(problem, strategy) returns a solution or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
```

Search Strategies

- A **search strategy** is defined by picking the **order of node expansion** e.g. Breadth First Search
- Strategies are evaluated along the following dimensions:
 1. **Completeness**
 - Guarantees finding a solution whenever one exists
 2. **Optimality/Admissibility**
 - If a solution is found, is it **guaranteed to be an optimal** one? For example, is it the one with minimum cost?
 3. **Time Complexity**
 - How long (worst or average case) does it take to find a solution? Usually measured in terms of the **number of nodes expanded**
 4. **Space Complexity**
 - How much space is used by the algorithm? Usually measured in terms of the **maximum size that the "OPEN" list** becomes during the search

Evaluating Search Strategies

- Time and space complexity are measured in terms of
 - b : maximum branching factor of the search tree
 - d : depth of the least-cost solution
 - m : maximum depth of the state space (may be ∞)

The End